

Q1. PCA for Breast Cancer Diagnosis (Real Data) (4×1 = 4 marks)

💡 Context

The Wisconsin Breast Cancer dataset contains 30 features per biopsy.

🎯 Goal

Use Principal Component Analysis (PCA) to denoise the data and visualize class separability.

📁 Data Source

Use either of the following:

- `sklearn.datasets.load_breast_cancer()`
 - Kaggle copy: [Breast Cancer Wisconsin Dataset](#)
-

```
In [85]: import numpy as np
import pandas as pd
from sklearn import datasets
data = datasets.load_breast_cancer()
```

```
In [86]: data.keys()
```

```
Out[86]: dict_keys(['data', 'target', 'frame', 'target_names', 'DESCR', 'feature_names', 'filename', 'data_module'])
```

```
In [87]: data.data.shape, data.target.shape
```

```
Out[87]: ((569, 30), (569,))
```

```
In [88]: df = pd.DataFrame(data.data, columns=data.feature_names)
df['target']=data.target
print(df.shape)
print(df.head())
```

```
(569, 31)
   mean radius  mean texture  mean perimeter  mean area  mean smoothness \
0      17.99      10.38       122.80     1001.0      0.11840
1      20.57      17.77       132.90     1326.0      0.08474
2      19.69      21.25       130.00     1203.0      0.10960
3      11.42      20.38        77.58      386.1      0.14250
4      20.29      14.34       135.10     1297.0      0.10030

   mean compactness  mean concavity  mean concave points  mean symmetry \
0      0.27760      0.3001       0.14710      0.2419
1      0.07864      0.0869       0.07017      0.1812
2      0.15990      0.1974       0.12790      0.2069
3      0.28390      0.2414       0.10520      0.2597
4      0.13280      0.1980       0.10430      0.1809

   mean fractal dimension  ...  worst texture  worst perimeter  worst area \
0      0.07871     ...      17.33       184.60     2019.0
1      0.05667     ...      23.41       158.80     1956.0
2      0.05999     ...      25.53       152.50     1709.0
3      0.09744     ...      26.50        98.87     567.7
4      0.05883     ...      16.67       152.20     1575.0

   worst smoothness  worst compactness  worst concavity  worst concave points \
0      0.1622       0.6656       0.7119      0.2654
1      0.1238       0.1866       0.2416      0.1860
2      0.1444       0.4245       0.4504      0.2430
3      0.2098       0.8663       0.6869      0.2575
4      0.1374       0.2050       0.4000      0.1625

   worst symmetry  worst fractal dimension  target
0      0.4601       0.11890      0
1      0.2750       0.08902      0
2      0.3613       0.08758      0
3      0.6638       0.17300      0
4      0.2364       0.07678      0
```

[5 rows x 31 columns]

Tasks

a) Standardize & Covariance

- Standardize all 30 features to have zero mean and unit variance.
- Compute the 30×30 sample covariance matrix (S).
- Report the Frobenius norm ($|S|_F$) as a numeric value.

```
In [89]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_std = scaler.fit_transform(data.data)
df2 = pd.DataFrame(X_std, columns=data.feature_names)
pd.concat([df2.mean(), df2.std()], axis=1, keys=['mean', 'std'])
```

Out[89]:

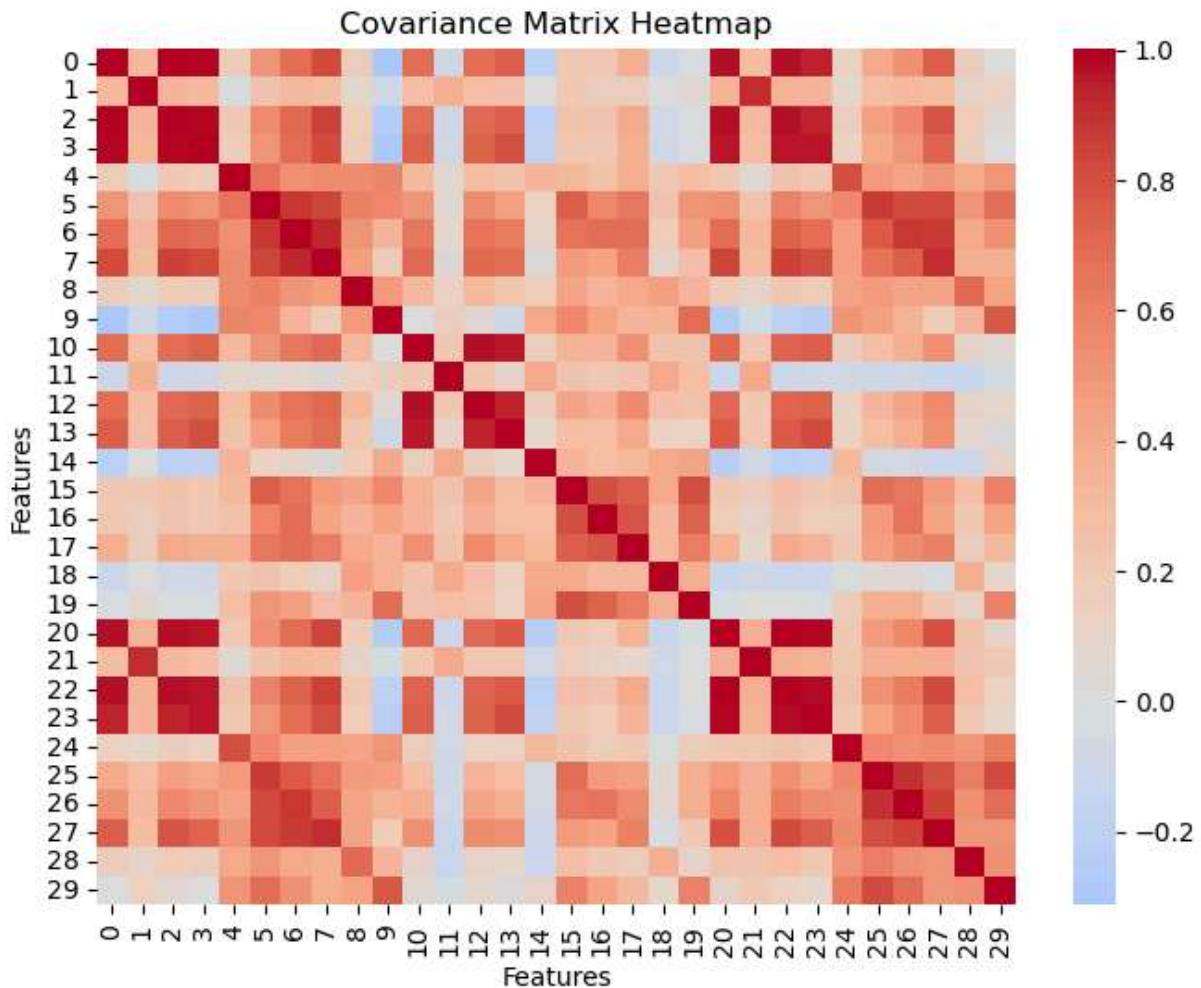
	mean	std
mean radius	-3.162867e-15	1.00088
mean texture	-6.530609e-15	1.00088
mean perimeter	-7.078891e-16	1.00088
mean area	-8.799835e-16	1.00088
mean smoothness	6.132177e-15	1.00088
mean compactness	-1.120369e-15	1.00088
mean concavity	-4.421380e-16	1.00088
mean concave points	9.732500e-16	1.00088
mean symmetry	-1.971670e-15	1.00088
mean fractal dimension	-1.453631e-15	1.00088
radius error	-9.076415e-16	1.00088
texture error	-8.853492e-16	1.00088
perimeter error	1.773674e-15	1.00088
area error	-8.291551e-16	1.00088
smoothness error	-7.541809e-16	1.00088
compactness error	-3.921877e-16	1.00088
concavity error	7.917900e-16	1.00088
concave points error	-2.739461e-16	1.00088
symmetry error	-3.108234e-16	1.00088
fractal dimension error	-3.366766e-16	1.00088
worst radius	-2.333224e-15	1.00088
worst texture	1.763674e-15	1.00088
worst perimeter	-1.198026e-15	1.00088
worst area	5.049661e-16	1.00088
worst smoothness	-5.213170e-15	1.00088
worst compactness	-2.174788e-15	1.00088
worst concavity	6.856456e-16	1.00088
worst concave points	-1.412656e-16	1.00088
worst symmetry	-2.289567e-15	1.00088
worst fractal dimension	2.575171e-15	1.00088

```
In [90]: df2.mean().mean(), df2.std().mean()
```

```
Out[90]: (np.float64(-6.369603796435407e-16), np.float64(1.000879894582902))
```

```
In [91]: # covariance matrix
S = np.cov(df2, rowvar=False)
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# printing heat map of covariance matrix
plt.figure(figsize=(8, 6))
sns.heatmap(S, annot=False, cmap='coolwarm', center=0)
plt.title("Covariance Matrix Heatmap")
plt.xlabel("Features")
plt.ylabel("Features")
plt.show()
```



```
In [92]: # Frobenius norm
fro_norm = np.linalg.norm(S, 'fro')
print("Frobenius norm of covariance matrix:", fro_norm)
```

```
Frobenius norm of covariance matrix: 15.062350986709806
```

b) Top-k Eigenpairs

- Set ($k = 2$).
- Use `sklearn.decomposition.PCA` to compute:
 - (i) Top two eigenvalues (λ_1, λ_2)
 - (ii) Explained variance ratio (%) of each
 - (iii) Cumulative explained variance (%)

In [93]:

```
# Eigvne value calculation as a reference, eigen values are printed in descending order
eigenvalues, eigenvectors = np.linalg.eig(S)
np.set_printoptions(precision=4, suppress=True)
idx = np.argsort(eigenvalues)[::-1]
print(eigenvalues[idx])
```

[13.305 5.7014 2.8229 1.9841 1.6516 1.2095 0.6764 0.4775 0.4176
 0.3513 0.2944 0.2616 0.2418 0.1573 0.0943 0.08 0.0595 0.0527
 0.0496 0.0312 0.03 0.0275 0.0244 0.0181 0.0155 0.0082 0.0069
 0.0016 0.0008 0.0001]

In [94]:

```
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
#df2 is the normalized breast cancer data (mean=0, sd = 1)
pca.fit(df2)
eigenvalues = pca.explained_variance_
print("Top 2 eigen values: ", eigenvalues)
explained_variance_ratio = pca.explained_variance_ratio_ * 100
print("Explained variance ratio: ", explained_variance_ratio)
cumulative_variance = explained_variance_ratio.cumsum()
print("Cumulative variance ratio: ", cumulative_variance)
```

Top 2 eigen values: [13.305 5.7014]
 Explained variance ratio: [44.272 18.9712]
 Cumulative variance ratio: [44.272 63.2432]

First 2 eigen values are covering up to ~82% of variance in the data

In [95]:

```
help(pca.fit)
```

Help on method fit in module sklearn.decomposition._pca:

```
fit(X, y=None) method of sklearn.decomposition._pca.PCA instance
  Fit the model with X.

  Parameters
  -----
  X : {array-like, sparse matrix} of shape (n_samples, n_features)
    Training data, where `n_samples` is the number of samples
    and `n_features` is the number of features.

  y : Ignored
    Ignored.

  Returns
  -----
  self : object
    Returns the instance itself.
```

c) 2D Projection

- Project all samples onto the top-2 principal components.
- Plot a labeled 2D scatter plot (malignant vs. benign).
- Report the class means in this 2D space.

```
In [96]: # projecting samples onto top-2 principal components
X_pca = pca.fit_transform(df2)
df_pca = pd.DataFrame(X_pca, columns=['PC1','PC2'])
df_pca['target'] = data.target
print(df_pca.head())
```

	PC1	PC2	target
0	9.192837	1.948583	0
1	2.387802	-3.768172	0
2	5.733896	-1.075174	0
3	7.122953	10.275589	0
4	3.935302	-1.948072	0

```
In [97]: X_pca.shape
```

```
Out[97]: (569, 2)
```

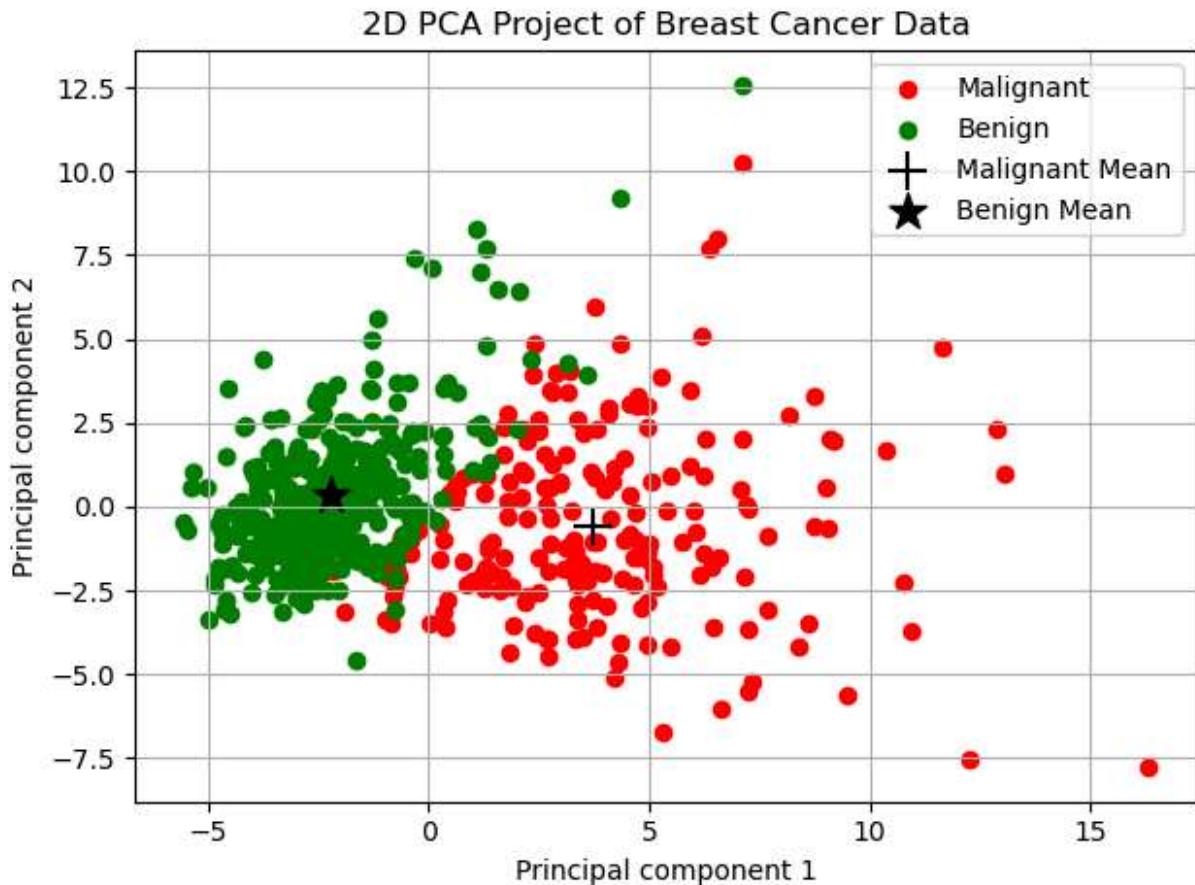
```
In [98]: malignant_features = df_pca[df_pca['target']==0][['PC1','PC2']]
benign_features = df_pca[df_pca['target']==1][['PC1','PC2']]
malignant_mean = malignant_features.mean().values
benign_mean = benign_features.mean().values
print("Malignant mean (PC1, PC2):", malignant_mean)
print("Benign mean (PC1, PC2):", benign_mean)

plt.scatter(malignant_features['PC1'], malignant_features['PC2'], c='red', label='Malignant')
plt.scatter(benign_features['PC1'], benign_features['PC2'], c='green', label='Benign')

plt.scatter(*malignant_mean, c='k', marker='+', s=200, label='Malignant Mean')
plt.scatter(*benign_mean, c='k', marker='*', s=200, label='Benign Mean')
plt.xlabel('Principal component 1')
plt.ylabel('Principal component 2')
plt.title('2D PCA Project of Breast Cancer Data')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

Malignant mean (PC1, PC2): [3.7148 -0.5831]

Benign mean (PC1, PC2): [-2.206 0.3463]



d) Reconstruction Error

- Inverse-transform the 2D projection back to the original 30D space.
- Report the average per-sample reconstruction Mean Squared Error (MSE) as a numeric value.

```
In [99]: X_recon = pca.inverse_transform(X_pca)
X_recon.shape
```

```
Out[99]: (569, 30)
```

```
In [100...]: from sklearn.metrics import mean_squared_error
mse = mean_squared_error(X_std, X_recon)
print(f"Average per-sample reconstruction MSE: {mse:.4f}")
```

```
Average per-sample reconstruction MSE: 0.3676
```

```
=====
```

```
In [ ]:
```